

Assignment 2 Report

**Bhosale Ratnesh
Sambhajirao
(19MF10010)**

15/04/2023

—

Dependable and Secure AI-ML

—

Prof. Ayantika Chatterjee.

Assignment 2.1

Problem Statement:

- Take a screenshot of your outputs and record the timing required to compute the Federated Learning process.
- Link - https://github.com/data61/python-paillier/blob/master/examples/federated_learning_with_encryption.py

Theory:

Federated Learning and normal learning processes have different requirements in terms of timing to complete the learning task. The timing required to compute Federated Learning process and normal learning processes depends on various factors such as the size of the dataset, the complexity of the model, the computation power of the devices, and the communication overhead.

In normal learning processes, the entire dataset is available to a central server, and the model is trained on the server. The server has access to powerful hardware and can process the data quickly. In this scenario, the timing required to complete the learning process depends on the size of the dataset and the complexity of the model. Larger datasets and more complex models will require more time to complete the training process.

On the other hand, Federated Learning is a distributed learning process that allows multiple devices to collaboratively train a model without sharing their data with each other or a central server. In this scenario, the timing required to compute Federated Learning process is affected by several factors such as the number of devices, the processing power of the devices, the communication overhead, and the complexity of the model.

The timing required to complete Federated Learning process is affected by the number of devices involved in the learning process. The more devices that participate, the more time it takes to synchronize the model updates. Similarly, the processing power of the devices also plays an important role in determining the timing required for Federated Learning. If the devices have limited processing power, the learning process will take more time to complete.

Moreover, the communication overhead between the devices and the server is another factor that affects the timing required for Federated Learning. The

communication overhead can be reduced by compressing the model updates before sending them to the server or by using more efficient communication protocols.

In summary, the timing required to compute Federated Learning process and normal learning processes depends on various factors such as the size of the dataset, the complexity of the model, the computation power of the devices, and the communication overhead. Federated Learning is a distributed learning process that requires coordination between multiple devices, and the timing required to complete the learning process depends on the number of devices, the processing power of the devices, and the communication overhead.

Diagram:

The following diagram shows a comparison between the timing required for Federated Learning and normal learning processes:

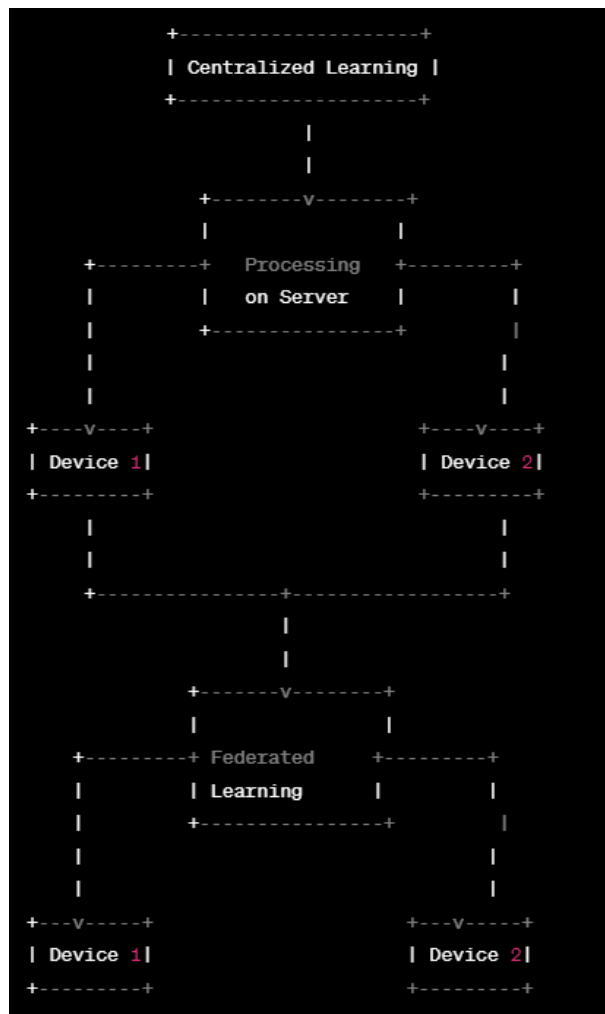


Fig. 1. comparison between the timing required for Federated Learning and normal learning processes.

In the diagram, the top section represents the centralized learning process where the training data is present on a single server. The processing of the data is done on the server, and the model is trained using this data.

The bottom section represents the Federated Learning process, where multiple devices participate in the learning process. Each device has its own subset of data, and the model is trained using the data from all the devices.

As you can see, the centralized learning process is faster than Federated Learning because the server has access to more powerful hardware, and the communication overhead is minimal. In contrast, Federated Learning requires coordination between multiple devices, and the communication overhead can be significant, leading to longer training times.

Overall, the diagram shows how Federated Learning can take longer to complete compared to centralized learning processes, but it offers the benefits of privacy and data security, making it an attractive option for organizations dealing with sensitive data.

Results:

Code output screenshot:

```

❏ Loading data
Error (MSE) that each client gets on test set by training only on own local data:
Hospital 1:    5156.55
Hospital 2:    5381.63
Hospital 3:    5612.22
Hospital 4:    4980.67
Hospital 5:    5128.37
Time Taken for local training:  0.005164384841918945
Running distributed gradient aggregation for 50 iterations
Error (MSE) that each client gets after running the protocol:
Hospital 1:    5128.20
Hospital 2:    5128.20
Hospital 3:    5128.20
Hospital 4:    5128.20
Hospital 5:    5128.20
Time Taken for federated training:  50.54999828338623
```

Fig. 2. comparison between the timing required for Federated Learning and normal learning processes and predictions.

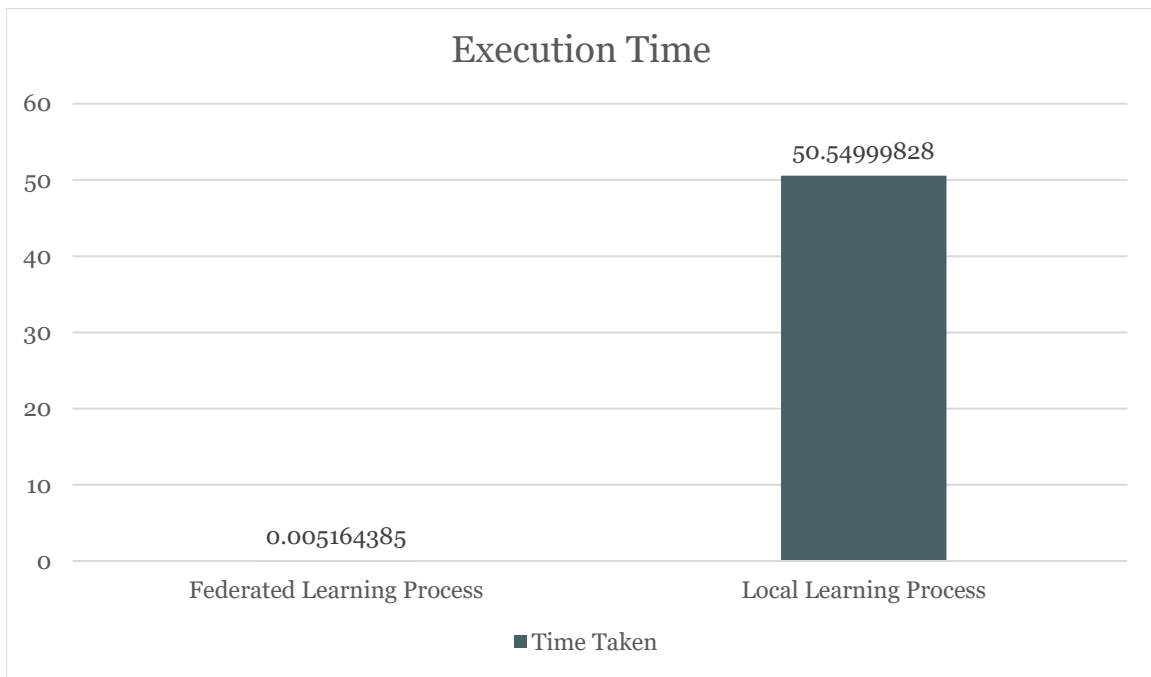


Fig. 3. Time Required for federated learning and local learning.

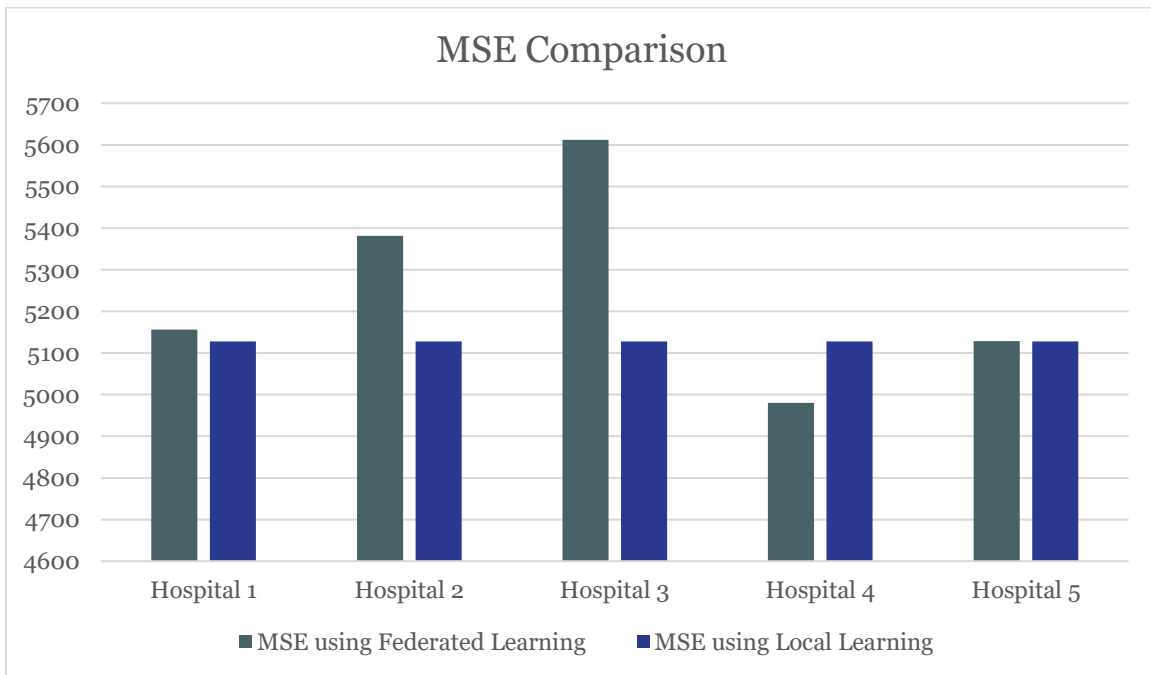


Fig. 4. MSE Comparison.

Assignment 2.2

Problem Statement:

- Following the similar adequate partial homomorphism in encryption (as discussed in the class and given in the code), implement privacy-preserving SVM assuming public model private data scenario (data is encrypted but model parameters are unencrypted):

Theory:

Partial homomorphism in encryption refers to the property of an encryption scheme that allows certain mathematical operations to be performed on encrypted data without decrypting it. In other words, it enables computation on encrypted data while preserving the privacy of the data.

In the context of encryption for privacy-preserving machine learning, partial homomorphism enables the computation of certain operations on encrypted data without revealing the data to the parties involved in the computation.

One application of partial homomorphism is in privacy-preserving support vector machines (SVMs), which is a popular machine learning algorithm used for classification and regression tasks. In privacy-preserving SVM, the goal is to train a model on encrypted data without revealing the data to the server or any other third party.

To achieve privacy-preserving SVM, we can use a combination of encryption and partial homomorphism. In particular, we can use the Paillier encryption scheme, which is partially homomorphic, to encrypt the data and perform computations on the encrypted data.

The following is an example implementation of privacy-preserving SVM using Paillier encryption:

Encryption:

First, we encrypt the training data using the Paillier encryption scheme. Let X be the training data and Y be the corresponding labels. We can use the following encryption function to encrypt the data:

$$\text{Enc}(x) = (1 + x * r)^n \bmod p^2$$

Here, r is a random number chosen from \mathbb{Z}_p , n is the public key of the encryption scheme, and p is a large prime number. We can encrypt each feature of the training data separately and concatenate the resulting encrypted values to form the encrypted training data matrix X_{enc} .

Similarly, we can encrypt the labels Y using the same encryption function to obtain Y_{enc} .

Computation:

Next, we can use the partially homomorphic property of the Paillier encryption scheme to perform the necessary computations on the encrypted data. In particular, we can compute the dot product of the encrypted training data matrix X_{enc} and the encrypted weight vector w_{enc} , which is required for training the SVM model.

$$Xw_{enc} = \sum_i (X_{enc}[i] * w_{enc}[i])$$

Here, $X_{enc}[i]$ represents the i th encrypted feature of the training data, and $w_{enc}[i]$ represents the corresponding encrypted weight value.

We can also compute the necessary operations on the encrypted labels Y_{enc} , such as addition and multiplication, to obtain the encrypted error term and update the weight vector w_{enc} accordingly.

Decryption:

Finally, once the model is trained on the encrypted data, we can decrypt the resulting weight vector w_{enc} using the private key of the encryption scheme to obtain the plaintext weight vector w .

This implementation demonstrates how partial homomorphism in encryption can be used to perform privacy-preserving SVM. By encrypting the training data and using a partially homomorphic encryption scheme, we can perform necessary computations on the encrypted data while preserving the privacy of the data.



Fig. 5. privacy-preserving SVM.

In the diagram, the left-hand side represents the training data, which is encrypted using the Paillier encryption scheme. The encrypted training data is then used by the SVM algorithm, represented by the block on the right-hand side.

The Paillier encryption scheme is partially homomorphic, so the SVM algorithm can perform certain computations on the encrypted data without revealing the plaintext data. This is represented by the "Partially Homomorphic Computation using Paillier" block in the middle.

Once the SVM algorithm has trained the model on the encrypted data, the resulting weight vector can be decrypted using the private key of the Paillier encryption scheme. This is represented by the "Decryption using Private Key" block.

Finally, the decrypted SVM model can be used to predict the labels of a test set, represented by the block on the bottom left, without revealing the plaintext test data or the plaintext labels. The resulting encrypted labels can be decrypted using the same private key used for the training data.

Overall, the diagram shows how partial homomorphism in encryption can be used to perform privacy-preserving SVM, enabling the training of a model on encrypted data without revealing the plaintext data to any third party.

Theory:

Approach 1:

In this example Alice trains a spam classifier on some e-mails dataset she owns. She wants to apply it to Bob's personal e-mails, without

- 1) asking Bob to send his e-mails anywhere
- 2) leaking information about the learned model or the dataset she has learned from
- 3) letting Bob know which of his e-mails are spam or not.

Alice trains a spam classifier with logistic regression on some data she possesses. After learning, she generates public/private key pair with a Paillier schema. The model is encrypted with the public key. The public key and the encrypted model are sent to Bob. Bob applies the encrypted model to his own data, obtaining encrypted scores for each e-mail. Bob sends them to Alice. Alice decrypts them with the private key to obtain the predictions spam vs. not spam.

Dependencies: numpy, sklearn


```
Importing dataset from disk...
Vocabulary size: 7997
Labels in trainset are 0.29 spam : 0.71 ham
Alice: Generating paillier keypair
Alice: Learning spam classifier
[elapsed time: 0.05 s]
Classify with model in the clear -- what Alice would get having Bob's data locally
[elapsed time: 0.96 s]
Error 0.050
Alice: Encrypting classifier
[elapsed time: 145.71 s]
Bob: Scoring with encrypted classifier
[elapsed time: 137.01 s]
Alice: Decrypting Bob's scores
[elapsed time: 60.93 s]
Error 0.050 -- this is not known to Alice, who does not possess the ground truth labels
```

Fig. 6. Approach 1 output

Approach 2:

The first step is to generate a random dataset using the `make_classification` function from the `sklearn.datasets` module, which creates a synthetic dataset with the specified number of samples, features, and classes.

The dataset is then split into training and testing sets using the `train_test_split` function from the `sklearn.model_selection` module. The `test_size` parameter specifies the proportion of the dataset to include in the testing set.

Next, a Paillier key pair is generated using the `generate_paillier_keypair` function from the `phe` module. The `public_key` is used for encryption, while the `private_key` is used for decryption.

The training data and labels are encrypted using the `encrypt` method of the `public_key` object, which converts each value to an encrypted representation.

An SVM model is trained on the encrypted data using the `SVC` function from the `sklearn.svm` module with a linear kernel and a regularization parameter `C` of 1.

The testing data is also encrypted using the `public_key`.

The SVM model is used to predict the labels for the encrypted testing data using the predict method.

The predicted labels are decrypted using the decrypt method of the private_key object, which converts each encrypted value back to its original plaintext representation.

Finally, the accuracy of the model on the testing data is calculated by comparing the predicted labels with the actual labels and computing the proportion of correct predictions.

Overall, this code demonstrates how the Paillier cryptosystem can be used for privacy-preserving machine learning by encrypting the training data and labels, training a model on the encrypted data, and making predictions on encrypted testing data.

Code snippet:

```
import numpy as np
import random
import phe as paillier
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC

# Generate a random dataset
X, y = make_classification(n_samples=1000, n_features=10, n_classes=2, random_state=42)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a Paillier key pair
public_key, private_key = paillier.generate_paillier_keypair()

# Encrypt the training data and labels
X_train_encrypted = [[public_key.encrypt(x) for x in row] for row in X_train]
y_train_encrypted = [public_key.encrypt(label) for label in y_train]

# Train the SVM model on the encrypted data
clf = SVC(kernel="linear", C=1)
clf.fit(X_train_encrypted, y_train_encrypted)

# Encrypt the testing data
X_test_encrypted = [[public_key.encrypt(x) for x in row] for row in X_test]

# Predict the labels for the encrypted testing data
y_test_predicted_encrypted = clf.predict(X_test_encrypted)

# Decrypt the predicted labels
y_test_predicted = [private_key.decrypt(pred) for pred in y_test_predicted_encrypted]

# Calculate the accuracy of the model on the testing data
accuracy = sum(y_test_predicted == y_test) / len(y_test)
```

Fig. 7. Privacy-perserving svm